

Gesture Recognition Case Study:

Group: Harshul Agarwal and Mikhil Varshney

This is the write-up for all models built to recognize the gestures correctly

Below experiments were performed to train models for Gesture Recognition,

Following are the results:

Experiment Number	Model	Result	Decision + Explanation																				
1	<div>Conv3D [Vanilla]</div> <table><tr><td>Image-size</td><td>120X120</td></tr><tr><td>Layers</td><td>64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>SGD</td></tr><tr><td>Batch size</td><td>60</td></tr></table>	Image-size	120X120	Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	SGD	Batch size	60	Out of Memory Error	<ul style="list-style-type: none">Created a vanilla model with 2 Conv3D layers.Batch size was taken as 60 which gave Out of Memory Error.Next Step is to reduce the batch size to 40.						
Image-size	120X120																						
Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																						
Dense Layer	256 neurons																						
[Con3D filter size/ strides/ padding	(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same																						
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																						
Optimizer	SGD																						
Batch size	60																						
2	<div>Conv3D [Vanilla]</div> <table><tr><td>Image-size</td><td>120X120</td></tr><tr><td>Layers</td><td>64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>SGD</td></tr><tr><td>Batch size</td><td>40</td></tr></table>	Image-size	120X120	Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	SGD	Batch size	40	<table><tr><td>Trainable Parameters</td><td>88,701,701</td></tr><tr><td>Training Accuracy</td><td>87.30 %</td></tr><tr><td>Validation Accuracy</td><td>57.14 %</td></tr></table>	Trainable Parameters	88,701,701	Training Accuracy	87.30 %	Validation Accuracy	57.14 %	<ul style="list-style-type: none">This vanilla model with 2 conv3D layers performs well on training accuracy but have very low validation accuracy.Overfitting as well.Next step is to change the filter size to a small value (2,2,2) to have better validation accuracy
Image-size	120X120																						
Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																						
Dense Layer	256 neurons																						
[Con3D filter size/ strides/ padding	(3,3,3) / (1,1,1) / same (3,3,3) / (1,1,1) / same																						
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																						
Optimizer	SGD																						
Batch size	40																						
Trainable Parameters	88,701,701																						
Training Accuracy	87.30 %																						
Validation Accuracy	57.14 %																						

3	Conv3D [Vanilla] (changed filter size to (2,2,2), Otherwise same as model 1	<table><tr><td>Trainable Parameters</td><td>88,483,717</td></tr><tr><td>Training Accuracy</td><td>85.71 %</td></tr><tr><td>Validation Accuracy</td><td>75 %</td></tr></table>	Trainable Parameters	88,483,717	Training Accuracy	85.71 %	Validation Accuracy	75 %	<ul style="list-style-type: none">As expected smaller filter size improved both accuracies.Overfitting got reduced by a large extent.Next step is to increase filter size to (5,5,5) to validate that bigger filter will perform bad.												
Trainable Parameters	88,483,717																				
Training Accuracy	85.71 %																				
Validation Accuracy	75 %																				
4	Conv3D [Vanilla] (changed filter size to (5,5,5), Otherwise same as model 1	<table><tr><td>Trainable Parameters</td><td>89,523,333</td></tr><tr><td>Training Accuracy</td><td>28.71 %</td></tr><tr><td>Validation Accuracy</td><td>31 %</td></tr></table>	Trainable Parameters	89,523,333	Training Accuracy	28.71 %	Validation Accuracy	31 %	<ul style="list-style-type: none">As expected bigger filter size worsened both accuracies.Next step is to try for lower image size for reducing training time, reverting to filter size of 2												
Trainable Parameters	89,523,333																				
Training Accuracy	28.71 %																				
Validation Accuracy	31 %																				
5	Conv3D [Vanilla] (changed image size 80X80) <table><tr><td>Image-size</td><td>80X80</td></tr><tr><td>Layers</td><td>64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>SGD</td></tr></table>	Image-size	80X80	Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	SGD	<table><tr><td>Trainable Parameters</td><td>88,492,101</td></tr><tr><td>Training Accuracy</td><td>83.34 %</td></tr><tr><td>Validation Accuracy</td><td>71.43 %</td></tr></table>	Trainable Parameters	88,492,101	Training Accuracy	83.34 %	Validation Accuracy	71.43 %	<ul style="list-style-type: none">Smaller image size as expected performed worse than larger.Switch back to (120X120) size as model perform best on it.Next Step is to reduce overfitting by trying different optimized (Adam)
Image-size	80X80																				
Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																				
Dense Layer	256 neurons																				
[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same																				
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																				
Optimizer	SGD																				
Trainable Parameters	88,492,101																				
Training Accuracy	83.34 %																				
Validation Accuracy	71.43 %																				

6	<div>Conv3D [Vanilla] (changed optimizer to Adam)</div> <table><tr><td>Image-size</td><td>120X120</td></tr><tr><td>Layers</td><td>64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>Adam</td></tr></table>	Image-size	120X120	Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	Adam	<table><tr><td>Trainable Parameters</td><td>88,542,405</td></tr><tr><td>Training Accuracy</td><td>40.65 %</td></tr><tr><td>Validation Accuracy</td><td>28.57 %</td></tr></table>	Trainable Parameters	88,542,405	Training Accuracy	40.65 %	Validation Accuracy	28.57 %	<ul style="list-style-type: none">Performed worse than SGD optimized.Will continue with SGD optimized.Next Step:- Add batch Normalization layer and revert back to SGD optimizer
Image-size	120X120																				
Layers	64-128 neurons with max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																				
Dense Layer	256 neurons																				
[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same																				
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																				
Optimizer	Adam																				
Trainable Parameters	88,542,405																				
Training Accuracy	40.65 %																				
Validation Accuracy	28.57 %																				
7	<div>Conv3D [Vanilla + Batch Normalization]</div> <table><tr><td>Image-size</td><td>120X120</td></tr><tr><td>Layers</td><td>64-128 layer neurons with Batch Norm and max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>SGD</td></tr></table>	Image-size	120X120	Layers	64-128 layer neurons with Batch Norm and max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	SGD	<table><tr><td>Trainable Parameters</td><td>88,548,573</td></tr><tr><td>Training Accuracy</td><td>89.34 %</td></tr><tr><td>Validation Accuracy</td><td>75.43 %</td></tr></table>	Trainable Parameters	88,548,573	Training Accuracy	89.34 %	Validation Accuracy	75.43 %	<ul style="list-style-type: none">Batch normalization improved the model performance and accuracy slightly.Next step, try adding dropout (0.25) layers between 64-128 neurons layer of model.
Image-size	120X120																				
Layers	64-128 layer neurons with Batch Norm and max pooling -> Flatten() -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																				
Dense Layer	256 neurons																				
[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same																				
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																				
Optimizer	SGD																				
Trainable Parameters	88,548,573																				
Training Accuracy	89.34 %																				
Validation Accuracy	75.43 %																				
8	<div>Conv3D [Vanilla + 3rd layer + Conv3D Dropout between 2nd and 3rd layer]</div> <table><tr><td>Image-size</td><td>120X120</td></tr><tr><td>Layers</td><td>32 neurons with max pooling and batch Norm -> Dropout (0.25)-> 64 neurons with max pooling -> Dropout (0.25)-> 128 neurons with max pooling -> Flatten () -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)</td></tr><tr><td>Dense Layer</td><td>256 neurons</td></tr><tr><td>[Con3D filter size/ strides/ padding</td><td>(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same</td></tr><tr><td>Max pool size / strides / padding</td><td>(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid</td></tr><tr><td>Optimizer</td><td>SGD</td></tr></table>	Image-size	120X120	Layers	32 neurons with max pooling and batch Norm -> Dropout (0.25)-> 64 neurons with max pooling -> Dropout (0.25)-> 128 neurons with max pooling -> Flatten () -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)	Dense Layer	256 neurons	[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same	Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid	Optimizer	SGD	<table><tr><td>Trainable Parameters</td><td>92,436,531</td></tr><tr><td>Training Accuracy</td><td>87.36 %</td></tr><tr><td>Validation Accuracy</td><td>79.67 %</td></tr></table>	Trainable Parameters	92,436,531	Training Accuracy	87.36 %	Validation Accuracy	79.67 %	<ul style="list-style-type: none">Adding additional layer to improve training accuracy and Adding dropouts between 2nd and 3rd layer to reduce the overfitting, increased the model validation accuracy slightly
Image-size	120X120																				
Layers	32 neurons with max pooling and batch Norm -> Dropout (0.25)-> 64 neurons with max pooling -> Dropout (0.25)-> 128 neurons with max pooling -> Flatten () -> Dropout (0.50) -> Dense (256) -> Dropout (0.50) -> Dense(5,softmax)																				
Dense Layer	256 neurons																				
[Con3D filter size/ strides/ padding	(2,2,2) / (1,1,1) / same (2,2,2) / (1,1,1) / same																				
Max pool size / strides / padding	(2,2,2) / (2,2,2) / valid (2,2,2) / (2,2,2) / valid																				
Optimizer	SGD																				
Trainable Parameters	92,436,531																				
Training Accuracy	87.36 %																				
Validation Accuracy	79.67 %																				

Conv3D Conclusion

- Dropouts are helping in reducing the overfitting by reducing training accuracy.
- (2,2,2) filter size is performing the best.
- 2 Conv3D layers are sufficient for achieving high training accuracy. But the model is overfitting.
- 3 Conv3D layers (model 8 with 256 layers in 3rd layer) can achieve a good fit and we can improve validation accuracy by using more training data

9	CNN(Conv2D) + RNN(GRU) Conv3D+BatchNorm+MaxPool repeated thrice followed by single GRU layer + 1 Dense layers. Dropout after each of the dense layers and GRU layer and then connect to the final output layer.	Trainable Parameters: 2,429,285 Val Accuracy : 56.25% Training Accuracy: 98.47%	<ul style="list-style-type: none">▪ We tried to create a CNN network that can learn & extract features and then fed that into RNN to perform classification of the gesture▪ The Current Model with 3 Con2D layer each followed by Maxpooling, batch Normalization & then by Flatten & dense layers with dropouts overfits & achieves a validation accuracy of 56.25%▪ We can see that we need more dense CNN network so we will try transfer learning with predefined N/W
---	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	---------------------------------------------------------------------------------------------------------------	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

10	CNN(Resnet) Earlier layers as fixed (last 5 layers trainable) + RNN(GRU)	Trainable Parameters: 6,560,773 Val Accuracy : 37.5% Training Accuracy: 89.6%	<ul style="list-style-type: none"> ▪ Trainable parameters increase with RESNET. ▪ We only train the last 5 layers to keep params and training time low. ▪ The Validation accuracy as observed is low. ▪ As feature extraction is responsibility of CNN, not training the weights of entire Resnet does not result in good validation accuracy.
----	---------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------	------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

11	CNN(MobileNet)All parameters trainable)+ 2 RNN(GRU) layers having dropout layers GRU (32, return_sequences = True) -> Dropout (0.5) -> GRU (64) -> Dropout (0.25)	Trainable Parameters : 3,331,653 Val Accuracy : 89.25% Training Accuracy: 99.09%	<ul style="list-style-type: none"> ▪ Now lets use another Predefined network for transfer learning : MobileNet ▪ We will train all the weights as seen from previous models that works better. ▪ This is much better than Resnet. But too much fluctuation can be seen in later epochs and overfitting is also noticeable.
----	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------------------	-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------

12. Final Model	<p>CNN(MobileNet)All parameters trainable) + 2 RNN(GRU) layers having dropout layers + L2 regularization in dense layers</p> <p>GRU (32, return_sequences = True) -> Dropout (0.25) -> GRU (64) -> Dropout (0.25) + Add L2 regularization</p>	<p>Trainable Parameters : 3,331,653 Val Accuracy : 95.83% Training Accuracy: 99.41%</p>	<ul style="list-style-type: none"> ▪ Now adding: ▪ Kernel regularizer = L2(0.01) ▪ Activity Regularizer = L2(0.01) ▪ This is the best model with less over fitting and the accuracy curve is converging.
------------------------	---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------	--------------------------------------------------------------------------------------------------------------------------	----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------